CLAIM_BOUNDARY — EP-AGENT-001

CLAIM BOUNDARY — EP-AGENT-001

Version: v0.3.4-DRAFT (2026-10-16; pending-narrowing pointer added) — superseded: v0.1-DRAFT (2026-10-15) — Status: DESIGN ONLY. No results exist; no claim below is currently made.

Principle: **claims remain narrower than evidence.**

Pending narrowing (v0.3.3 pointer; not yet operative): if the PREREGISTRATION §18 amendment (DRAFT, NOT FILED) is approved and filed, the §1 claim refers only to the `/verify` decision path of the pinned ExecutionProof Controlled Reference Verifier (`v0.1.2-reference` @ `b03365d57880e896c35e72898f8335a1826afe74`, with the study policy and evidence registry hashes in `EXECUTIONPROOF_PIN.md`) and must carry the reference-implementation disclosure. It may not be broadened to any other ExecutionProof build, deployment, or version. The text below is otherwise unchanged.

---

1. Maximum allowed successful claim (verbatim; usable only if the classification is PASS)

> "In an RF-run internal study, the ExecutionProof `/verify` decision path was integrated with four selected agent frameworks (OpenAI Agents SDK 0.24.0, LangGraph 1.2.15, CrewAI 1.15.28, Microsoft Agent Framework 1.22.0) through thin single-agent adapters using one frozen Request Contract. Across five benign, policy-driven fixtures, all 20 framework × case cells returned the preregistered ALLOW / HOLD / DENY verdict with ProofRecord linkage, without modification of the ExecutionProof core decision engine."

For PARTIAL, the claim must name exactly which frameworks met which endpoints and which did not. For FAIL, the failure is reported first and in full. For INCONCLUSIVE, the invalidation condition is named and no portability claim is made.

2. Required qualifiers (every external or internal summary must carry all of these)

1. "RF-run internal study" — never "independent", "third-party", or "external".
2. The exact framework versions tested; no generalization to other versions.
3. "Single-agent tool-call path only"; delegation/handoff features not used.
4. "Five benign, policy-driven fixtures"; "not an adversarial evaluation".
5. "Local/sandbox ExecutionProof instance with mock downstream targets"; not production.
6. Latency figures labelled "RF internal measurement" with host spec.
7. The non-interference statement: "It makes no claim about adversarial robustness, multi-agent authorization, or independent reproducibility, which are the subject of independent university projects."

3. Prohibited claims

#	Prohibited claim or wording	Reason
1	"ExecutionProof works with any / all agent frameworks"	Only four pinned versions tested
2		Overgeneralization

#	Prohibited claim or wording	Reason
	"Framework-agnostic" or "universally portable" without the tested-version qualifier	
3	"Secure", "attack-resistant", "tamper-proof", "bypass-proof", or "robust" on the basis of this study	No adversarial testing (UD / UW P1)
4	"Prevents replay", "prevents mutation", "no false-ALLOW" on the basis of this study	Reserved; not tested
5	Any claim about multi-agent, delegated, or agent-to-agent authorization	Reserved (UW P2); not tested
6	"Independently validated", "independently reproduced", "reproduced by", "externally verified"	Reserved (UW P4); RF-run
7	"Conformant", "RF-100 conformant", "certified", "compliant"	Reserved (UW P4); RF cannot self-certify
8	Calling any university "a validator" or implying university endorsement	Rule 7
9	"Production-ready" or "enterprise-proven" on the basis of this study	Sandbox only
10	"Zero overhead", "negligible latency", or latency comparisons as general performance claims	Internal, host-specific measurement
11	"Novel" / "first" without a cited literature review supporting it	Novelty not established by this study
12	Framing results as answering, pre-empting, or anticipating UD or UW findings	Rules 3–4
13	Any claim before preregistration SHA lock and authorized execution	No results exist
14	Re-describing a FAIL or PARTIAL as PASS after remediation	Rule 11; FAIL and fix are two facts

4. Scientific vs. commercial separation rule

1. The scientific record (preregistration, results, classification) is written first and stands alone.
2. The commercial question ("Can an implementation partner retain its preferred agent framework while ExecutionProof remains an independent pre-execution verification boundary?") is answered, if at all, in a separate document that cites the scientific record and inherits every qualifier in §2.
3. Commercial material may not upgrade the classification, drop qualifiers, or round PARTIAL up to PASS.
4. "Independent" in the CQ refers to architectural separation of the verification boundary from the agent framework — never to independent validation. Commercial copy must make that distinction explicit.
5. No result is published or archived (including to Zenodo) without Derek Hone's separate publication authorization.